

p8106_hw4_jsg2145

Jared Garfinkel

4/25/2020

Part a


```
library(lasso2)
```

```
## R Package to solve regression problems while imposing
##   an L1 constraint on the parameters. Based on S-plus Release 2.1
## Copyright (C) 1998, 1999
## Justin Lokhorst <jlokhors@stats.adelaide.edu.au>
## Berwin A. Turlach <bturlach@stats.adelaide.edu.au>
## Bill Venables <wvenable@stats.adelaide.edu.au>
##
## Copyright (C) 2002
## Martin Maechler <maechler@stat.math.ethz.ch>
```

```
data(Prostate)
```

```
x <- model.matrix(lpsa~.,Prostate)[,-1]
y <- Prostate$lpsa
```

```
tree1 <- rpart(lpsa~., Prostate, control = rpart.control(cp = .0001))
rpart.plot(tree1)
```



```
tree1$cptable
```


```
##          CP nsplit rel error    xerror    xstd
## 1 0.34710828    0 1.0000000 1.0159325 0.16219447
## 2 0.18464743    1 0.6528917 0.8818016 0.11403864
## 3 0.05931585    2 0.4682443 0.6807611 0.08252449
## 4 0.03475635    3 0.4089284 0.5613083 0.06425663
## 5 0.03460901    4 0.3741721 0.5914809 0.06606627
## 6 0.02156368    5 0.3395631 0.6152283 0.06669151
## 7 0.02146995    6 0.3179994 0.6180729 0.06603792
## 8 0.00010000    7 0.2965295 0.5579744 0.05938270
```

```
cpTable <- printcp(tree1)
```


```
##
## Regression tree:
## rpart(formula = lpsa ~ ., data = Prostate, control = rpart.control(cp = 1e-04))
##
## Variables actually used in tree construction:
## [1] lcavol lweight pgg45
##
## Root node error: 127.92/97 = 1.3187
##
## n= 97
##
```

##	CP	nsplit	rel error	xerror	xstd
## 1	0.347108	0	1.00000	1.01593	0.162194
## 2	0.184647	1	0.65289	0.88180	0.114039
## 3	0.059316	2	0.46824	0.68076	0.082524
## 4	0.034756	3	0.40893	0.56131	0.064257
## 5	0.034609	4	0.37417	0.59148	0.066066
## 6	0.021564	5	0.33956	0.61523	0.066692
## 7	0.021470	6	0.31800	0.61807	0.066038
## 8	0.000100	7	0.29653	0.55797	0.059383


```
plotcp(tree1)
```



```
minErr <- which.min(cpTable[,4])
# minimum cross-validation error
tree3 <- prune(tree1, cp = cpTable[minErr,1])
# 1SE rule
tree4 <- prune(tree1, cp = cpTable[cpTable[,4]<cpTable[minErr,4]+cpTable[minErr,5],1][1])
rpart.plot(tree3)
```



```
rpart.plot(tree4)
```




```
ctrl1 = trainControl(method = "repeatedcv", number = 10, repeats = 5)

set.seed(22)
tree_caret_cv = train(x, y, method = "rpart",
                      tuneGrid = data.frame(cp = seq(.001, 1, length = 1000)),
                      trControl = ctrl1)

tree_caret_cv$bestTune
```

```
##      cp
## 13 0.013
```


```
ggplot(tree_caret_cv, highlight = TRUE)
```



```
tree_caret_cv$finalModel$cptable
```

```
##          CP nsplit rel error
## 1 0.34710828      0 1.0000000
## 2 0.18464743      1 0.6528917
## 3 0.05931585      2 0.4682443
## 4 0.03475635      3 0.4089284
## 5 0.03460901      4 0.3741721
## 6 0.02156368      5 0.3395631
## 7 0.02146995      6 0.3179994
## 8 0.00000000      7 0.2965295
```

```
rpart.plot(tree_caret_cv$finalModel)
```



```

set.seed(22)
tree_caret_1se <- train(x, y,
  method = "rpart",
  tuneGrid = data.frame(cp = seq(.001, 1, length = 1000)),
  trControl = trainControl(method = "repeatedcv", number = 10, repeats = 5,
    selectionFunction = "oneSE"))

tree_caret_1se$bestTune

```

```

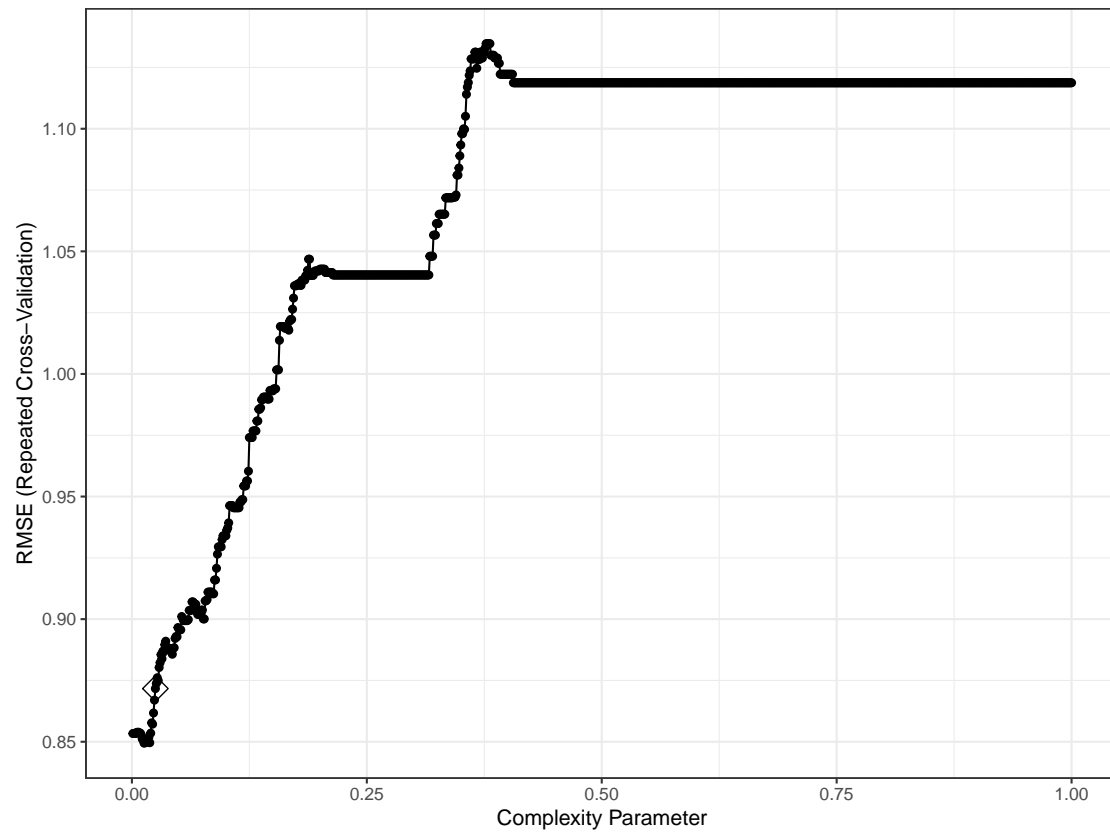
##          cp
## 25 0.025

```

```

ggplot(tree_caret_1se, highlight = TRUE) + theme_bw()

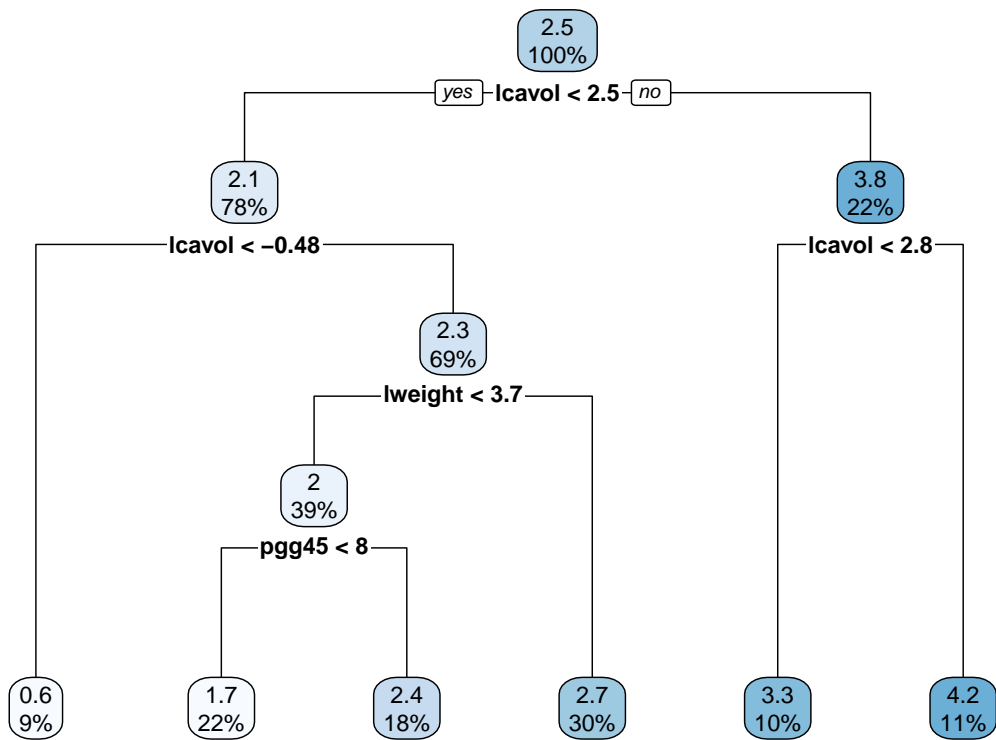
```



```
tree_caret_1se$finalModel$cptable
```

```
##          CP nsplit rel error
## 1 0.34710828     0 1.0000000
## 2 0.18464743     1 0.6528917
## 3 0.05931585     2 0.4682443
## 4 0.03475635     3 0.4089284
## 5 0.03460901     4 0.3741721
## 6 0.02500000     5 0.3395631
```

```
rpart.plot(tree_caret_1se$finalModel)
```





```

set.seed(22)
resamp <- resamples(list(minErr = tree_caret_cv,
                        oneSE = tree_caret_1se))

ggplot(resamp)

```




```
summary(resamp)
```

```
##
## Call:
## summary.resamples(object = resamp)
##
## Models: minErr, oneSE
## Number of resamples: 50
##
## MAE
##           Min.   1st Qu.   Median     Mean   3rd Qu.   Max. NA's
## minErr 0.4519217 0.6201402 0.7089144 0.7244433 0.8258116 1.010807    0
## oneSE  0.4712642 0.6527670 0.7582356 0.7366359 0.8261562 1.010807    0
##
## RMSE
##           Min.   1st Qu.   Median     Mean   3rd Qu.   Max. NA's
## minErr 0.5366273 0.7021520 0.8576688 0.8494137 0.9487340 1.171649    0
## oneSE  0.5739553 0.7572611 0.8939170 0.8716841 0.9600888 1.165925    0
##
## Rsquared
##           Min.   1st Qu.   Median     Mean   3rd Qu.   Max. NA's
## minErr 1.444536e-05 0.3509990 0.4734469 0.4871524 0.6473874 0.8345701    0
## oneSE  2.321338e-02 0.3337917 0.4426040 0.4561976 0.6124944 0.8016926    0
```

Part b

```
final_tree = rpart(formula = lpsa ~ ., data = Prostate, control = rpart.control(cp = 0.1))
rpart.plot(final_tree)
```

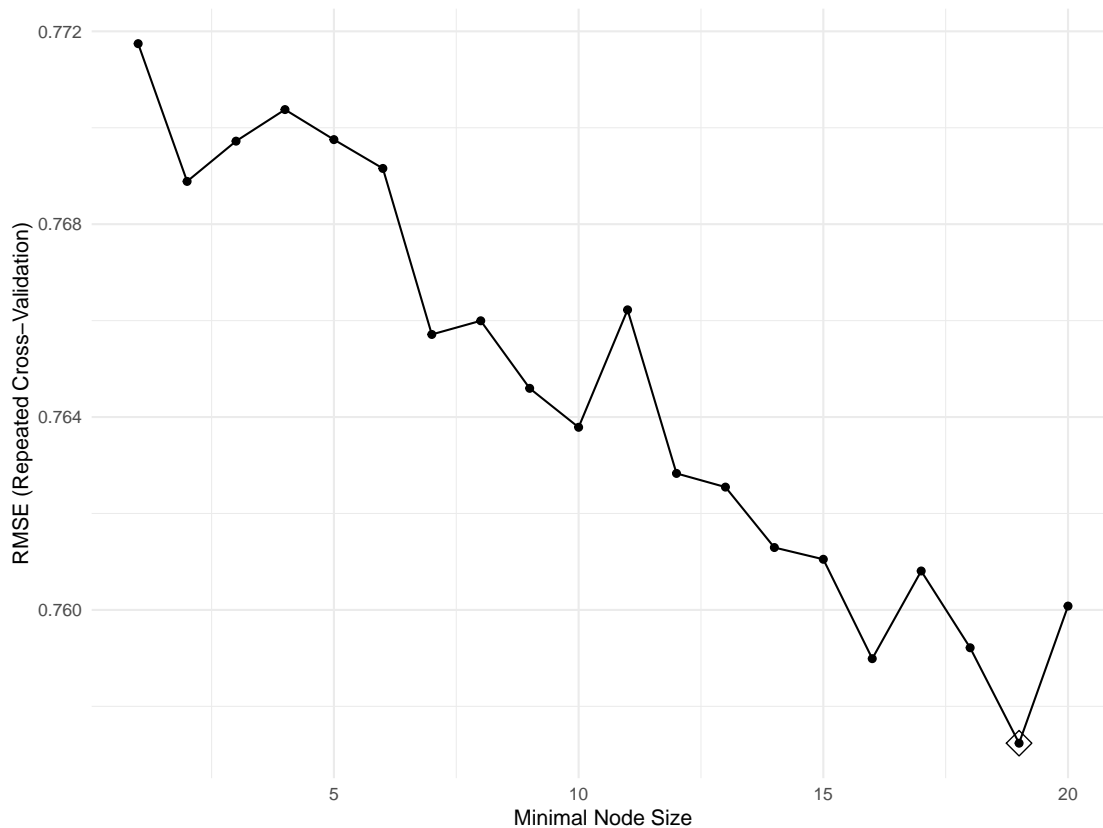


Part c

```
bagging_grid <- expand.grid(mtry = 8,
                             splitrule = "variance",
                             min.node.size = 1:20)

set.seed(22)
bagging_fit <- train(x, y,
                     method = "ranger",
                     tuneGrid = bagging_grid,
                     trControl = ctrl1,
                     importance = "impurity")


ggplot(bagging_fit, highlight = TRUE)
```



```
bagging_fit$results[which.min(bagging_fit$results[,5]),]
```


```
##      mtry splitrule min.node.size      RMSE Rsquared      MAE      RMSESD
## 20      8 variance          20 0.7600798 0.5993522 0.6307274 0.1594892
##      RsquaredSD      MAESD
## 20 0.1557015 0.1432667
```

```
barplot(sort(ranger::importance(bagging_fit$finalModel),
            decreasing = FALSE),
        las = 2,
        horiz = TRUE,
        cex.names = 0.7,
        col = colorRampPalette(colors = c("darkred",
                                           "white",
                                           "darkblue"))(19))
```



Part d


```
randfor_grid <- expand.grid(mtry = 1:7,  
                           splitrule = "variance",  
                           min.node.size = 1:15)  
  
set.seed(22)  
randfor_fit <- train(x, y,  
                    method = "ranger",  
                    tuneGrid = randfor_grid,  
                    trControl = ctrl1,  
                    importance = 'permutation')  
  
ggplot(randfor_fit, highlight = TRUE)
```



```
randfor_fit$results[which.min(randfor_fit$results[,5]),]
```


```
##      mtry splitrule min.node.size    RMSE Rsquared      MAE    RMSESD
## 14      1  variance          14 0.79687 0.5607409 0.6386306 0.1956834
##      RsquaredSD    MAESD
## 14 0.1515241 0.1477475
```

```
barplot(sort(ranger::importance(randfor_fit$finalModel), decreasing = FALSE),
        las = 2, horiz = TRUE, cex.names = 0.7,
        col = colorRampPalette(colors = c("darkred", "white", "darkblue"))(19))
```




Part e

```
gbm_grid <- expand.grid(  
  n.trees = seq(1, 5000, 100),  
  interaction.depth = 2:10,  
  shrinkage = c(0.001, 0.003, 0.005),  
  n.minobsinnode = 1)  
  
set.seed(22)  
gbm_fit <- train(x, y,  
  method = "gbm",  
  tuneGrid = gbm_grid,  
  trControl = ctrl1,  
  verbose = FALSE)  
  
ggplot(gbm_fit, highlight = T) + theme_bw()
```



```
summary(gbm_fit$finalModel, las = 2, cBars = 19, cex.names = 0.6)
```

```
##          var  rel.inf
## lcavol  lcavol 54.953514
## lweight lweight 17.796552
## svi      svi   7.396924
## pgg45    pgg45  5.933298
## lcp      lcp   5.624608
## age      age   4.381118
## lbph     lbph  2.419184
## gleason  gleason 1.494802
```

Part f


```
resamp2 = resamples(list(minErr = tree_caret_cv,
                          min_1se = tree_caret_1se,
                          randomForest = randfor_fit,
                          boosting = gbm_fit))

summary(resamp2)
```

```
##
## Call:
## summary.resamples(object = resamp2)
##
## Models: minErr, min_1se, randomForest, boosting
## Number of resamples: 50
```

```
##
## MAE
##           Min.    1st Qu.    Median      Mean    3rd Qu.      Max.
## minErr      0.4519217 0.6201402 0.7089144 0.7244433 0.8258116 1.0108066
## min_1se      0.4712642 0.6527670 0.7582356 0.7366359 0.8261562 1.0108066
## randomForest 0.3586373 0.5110769 0.6101846 0.6202786 0.6995581 0.9570076
## boosting     0.3662583 0.5171184 0.6064738 0.6124641 0.7035913 0.9653459
##           NA's
## minErr      0
## min_1se      0
## randomForest 0
## boosting     0
##
## RMSE
##           Min.    1st Qu.    Median      Mean    3rd Qu.      Max.
## minErr      0.5366273 0.7021520 0.8576688 0.8494137 0.9487340 1.171649
## min_1se      0.5739553 0.7572611 0.8939170 0.8716841 0.9600888 1.165925
## randomForest 0.5098127 0.6130050 0.7267068 0.7512219 0.8499004 1.118397
## boosting     0.4278317 0.6470290 0.7257828 0.7483030 0.8626489 1.051423
##           NA's
## minErr      0
## min_1se      0
## randomForest 0
## boosting     0
##
## Rsquared
##           Min.    1st Qu.    Median      Mean    3rd Qu.
## minErr      1.444536e-05 0.3509990 0.4734469 0.4871524 0.6473874
## min_1se      2.321338e-02 0.3337917 0.4426040 0.4561976 0.6124944
## randomForest 3.496955e-01 0.5102101 0.5926669 0.6082333 0.6851048
## boosting     3.563333e-01 0.4902781 0.6031274 0.6184048 0.7532405
##           Max. NA's
## minErr      0.8345701    0
## min_1se      0.8016926    0
## randomForest 0.8995859    0
## boosting     0.9214752    0
```

```
bwplot(resamp2, metric = "RMSE")
```



Problem 2

Problem 2a

```
data(OJ)
oj_data = OJ %>%
  janitor::clean_names()
# create a training set containing 800 obs
set.seed(22)
rowTrain = createDataPartition(y = oj_data$purchase,
                                p = 799/1070,
                                list = F)
train_data = oj_data[rowTrain, ]
test_data = oj_data[-rowTrain, ]
# check whether there is 800 obs
dim(train_data)
```

```
## [1] 800 18
```

```
x_train = train_data[,-1]
y_train = pull(train_data, purchase)
```


```
x_test = test_data[,-1]
y_test = pull(test_data, purchase)
```

```
ctrl2 <- trainControl(method = "repeatedcv",
                      summaryFunction = twoClassSummary,
                      classProbs = TRUE)

set.seed(22)

fit_oj_cv <- train(x_train, y_train,
                  method = "rpart",
                  tuneGrid = data.frame(cp = exp(seq(-20,-5, len = 20))),
                  trControl = ctrl2,
                  metric = "ROC")

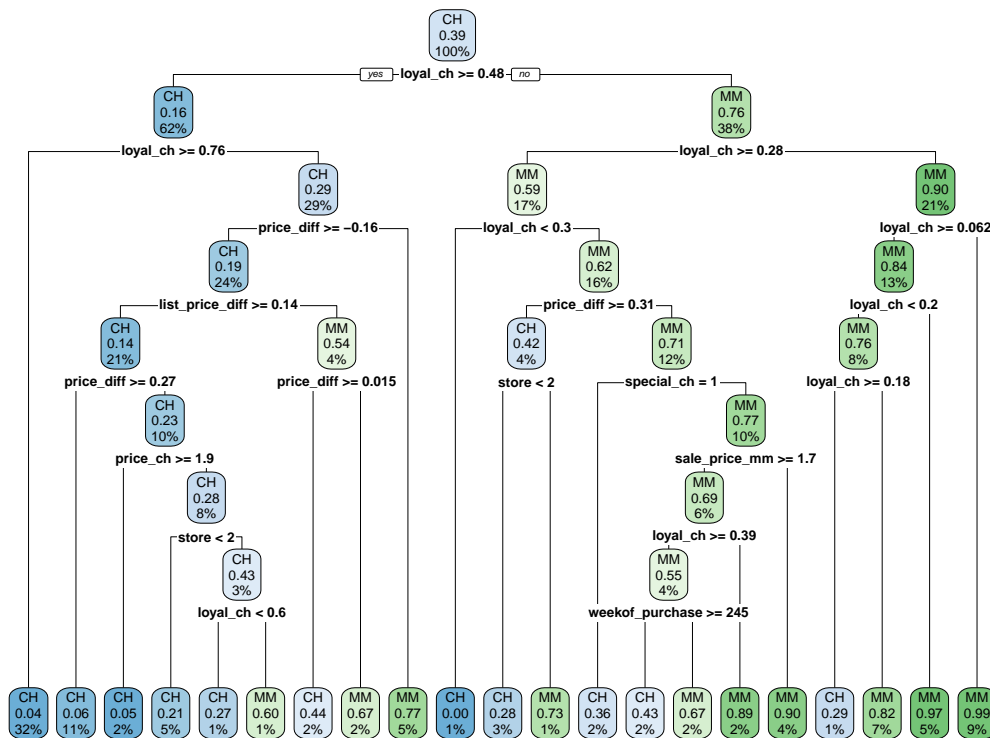
ggplot(fit_oj_cv, highlight = TRUE)
```



```
# optimal tree size is 17 with smallest CV error
fit_oj_cv$finalModel$cptable
```

```
##          CP nsplit rel error
## 1 0.519230769      0 1.0000000
## 2 0.033653846      1 0.4807692
## 3 0.011217949      3 0.4134615
## 4 0.006410256      8 0.3429487
## 5 0.003205128     10 0.3301282
```

```
# plot of tree
rpart.plot(fit_oj_cv$finalModel)
```

[illegible]

```
# test classification error rate
1 - mean(test_data$purchase == pred)
```

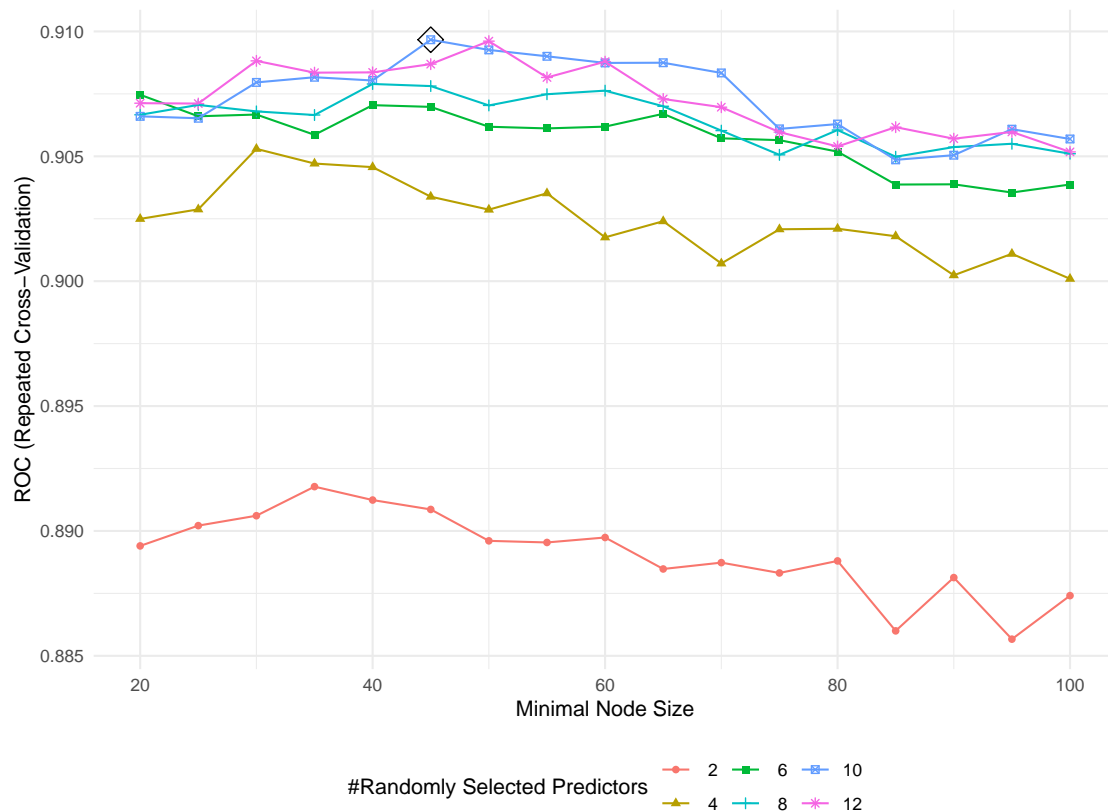
```
## [1] 0.2148148
```

Part b


```
rf.grid_oj <- expand.grid(mtry = seq(2,12,2),
  splitrule = "gini",
  min.node.size = seq(20,100,5))

set.seed(1)
rf.fit_oj <- train(purchase ~ ., train_data,
  method = "ranger",
  tuneGrid = rf.grid_oj,
  metric = "ROC",
  importance = "impurity",
  trControl = ctrl2)

# rf plot
ggplot(rf.fit_oj, highlight = TRUE)
```



```
# compare variable importance
barplot(sort(ranger::importance(rf.fit_oj$finalModel), decreasing = FALSE),
  las = 2, horiz = TRUE, cex.names = 0.7,
  col = colorRampPalette(colors = c("darkred", "white", "darkblue"))(19))
```



```
# predict on test data
pred2 = predict(rf.fit_oj, newdata = test_data,
               type = "raw");pred
```

```
## [1] CH CH CH CH CH CH CH CH MM CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH MM CH
## [24] CH MM CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH
## [47] CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH
## [70] CH CH CH CH CH CH CH MM MM CH MM MM MM MM MM MM MM MM MM MM MM CH CH CH CH
## [93] CH MM MM CH MM CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH
## [116] CH MM MM MM MM MM MM CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH
## [139] MM CH MM MM MM MM MM MM CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH
## [162] CH CH CH MM MM MM CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH
## [185] MM MM MM CH MM MM CH CH MM MM CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH
## [208] MM CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH CH
## [231] CH CH CH MM MM CH CH CH MM MM MM MM MM CH MM MM MM MM CH MM MM CH CH
## [254] CH CH CH CH MM CH MM MM CH CH CH CH CH CH CH CH CH CH CH CH CH CH
## Levels: CH MM
```

```
# test error rate
1 - mean(test_data$purchase == pred2)
```

```
## [1] 0.188889
```